NanoBook: Bare-Metal Session Report

Phases 3–5 published numbers and phase 6's re-checks · 24 October 2026 · C++23, GCC 14.3.0, Ubuntu 26.04.1 LTS on bare metal · Session run; decisions pending (§6)

1. Summary

Every number before this session came from WSL2, a virtual machine, and was marked "not for publishing" (README rule 20). This session ran the same tools on the same laptop, booted into Linux directly from a USB stick, with three physical cores set aside for the measurements. It does two jobs: it produces the **published numbers** for phases 3, 4 and 5, and it **re-checks phase 6's decisions** (optimise design §5, B1–B7), using rules written down before the data.

Status	**All runs done**, in one session on 24 October. Raw data in `docs/data/bare_metal/`. Three decisions wait for the user (§6)
Real Nasdaq day, one thread	**75 ns per message** (13.4 M messages/s); WSL2 at the end of phase 6: about 92 ns
Real day, slowest 1 in 10,000 messages (p99.99)	**0.65 µs**; WSL2: about 11 µs
Pipeline p99 at 1–4 M messages/s	**about 0.65 µs**; WSL2: 65–72 µs
Kept changes re-checked	E1b, E3, E6 **confirmed**; E2 **fails** its rule (revert)
Undecided change E9b	**passes** its rule: 7% faster, p50 24% better (re-add)
Moving window E7	15% faster, p99 now equal, but p99.99 2.5–6× worse

The one-sentence lesson: the virtual machine hid good news as well as bad: the pipeline's 65 µs tail was the VM, not the code, and E9b's gain was real but drowned in WSL2's noise.

2. The setup

2.1 The machine

CPU	Intel Core i7-14650HX (8 performance cores with two hyperthreads each, 8 efficiency cores; 24 logical CPUs)
OS	**Ubuntu 26.04.1 LTS**, live USB stick with a persistent partition, booted directly (no virtual machine)
Kernel	Linux 7.0.0-30-generic, boot options `isolcpus=8-13 nohz_full=8-13 rcu_nocbs=8-13`
Cores used	CPU 8 for one-thread runs; CPUs 8, 10, 12 for the pipeline (three different performance cores); their hyperthread twins 9, 11, 13 left idle
Frequency	`performance` governor, **turbo on** (CPU 8 ran at 5.2 GHz, read from the cycle counter)
Compiler	GCC 14.3.0, `gcc-release` preset: `-std=c++23 -O3 -DNDEBUG`
Code	Commit `943fe70` (clean tree, rebuilt before each session) unless the A/B says otherwise
Input	Nasdaq TotalView-ITCH 5.0, 30 January 2019, 11.2 GB, 368,366,634 messages, read-only from the Windows drive
Power	Plugged in, laptop idle

2.2 What the three boot options do

- `isolcpus=8-13`: Linux puts no other program on CPUs 8–13. Only our pinned threads run there.

- `nohz_full=8-13`: the timer tick (an interrupt every few milliseconds) stops on those CPUs while one thread runs, so the measured code isn't interrupted.
- `rcu_nocbs=8-13`: a kernel housekeeping job (RCU callbacks) moves to the other CPUs.

Together they answer B4: are the slow tails the code's fault or the machine's?

2.3 Every tool checked the setup itself

Each output starts with the environment the tool read from the system: `virtual machine no`, `isolated cores 8-13`, `governor performance`, `turbo on`. The "NOT FOR PUBLISHING" banner, printed whenever a hypervisor is found, is gone from every file.

---

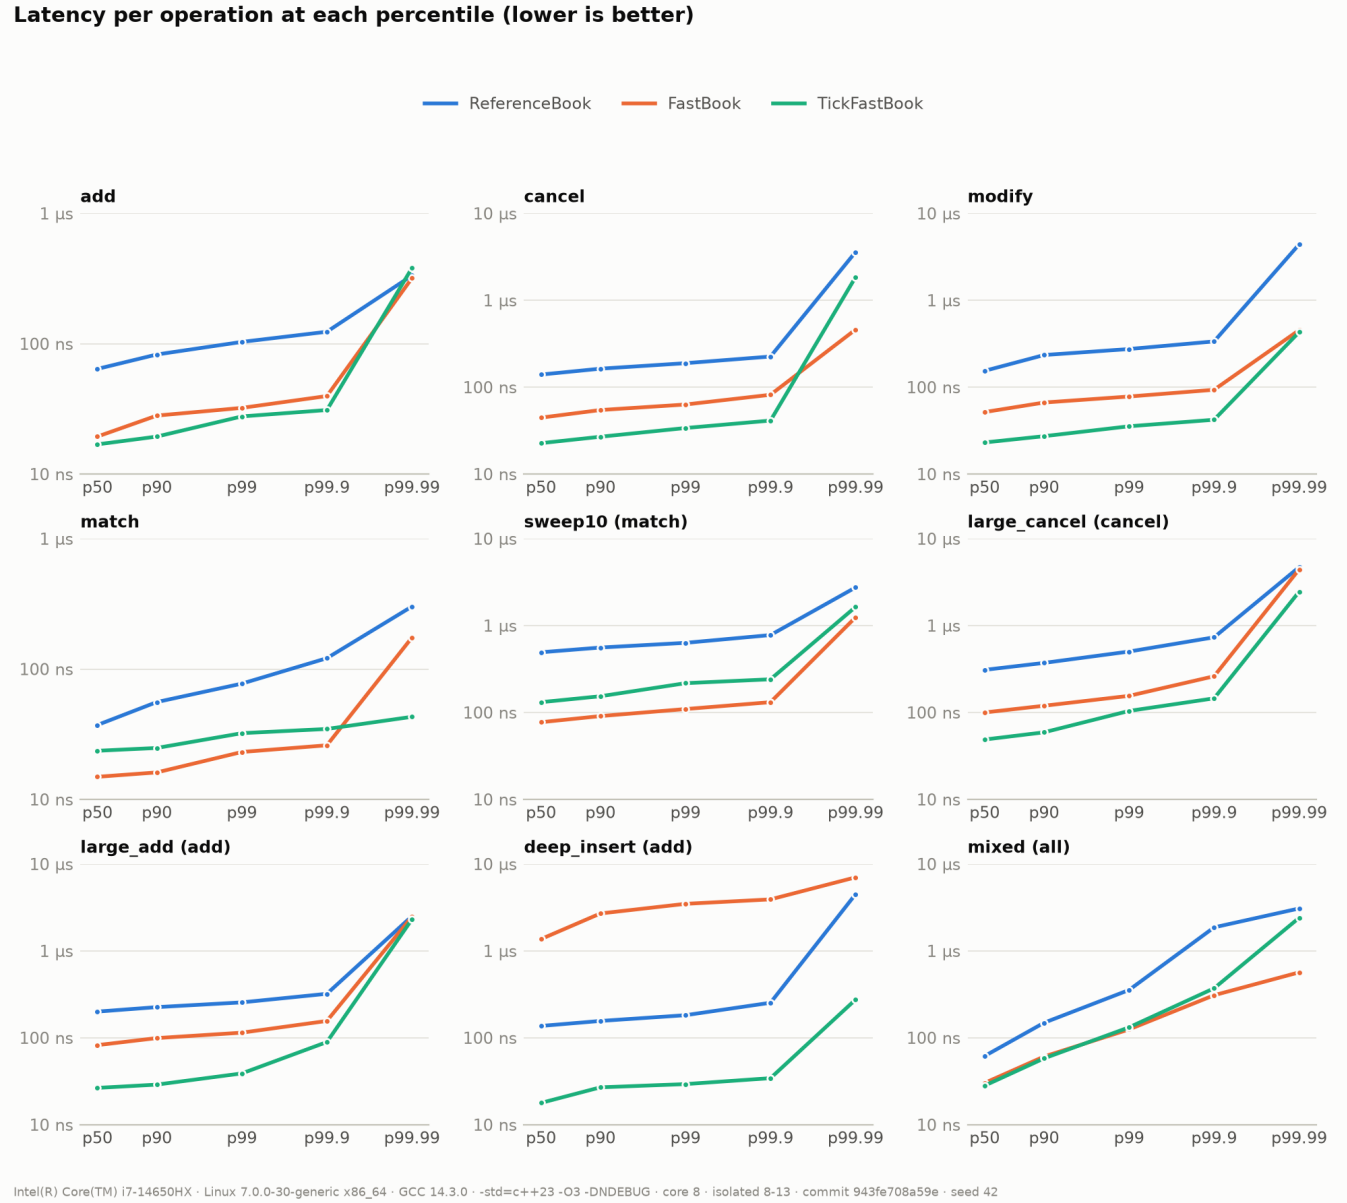
3. Published numbers

3.1 Phase 3: one operation at a time (synthetic books)

`scripts/run_latency_session.sh docs/data/bare_metal/phase3 --core 8`: 1,000,000 timed operations per scenario, two rounds. FastBook, round 1 / round 2, in ns:

Scenario	p50	p99	p99.99	WSL2 p99.99
add	19.4 / 20.3	32.2 / 30.2	318 / 299	2,229
cancel	44.6 / 44.2	63.2 / 62.4	461 / 137	2,433
modify	51.7 / 51.7	78.1 / 77.3	454 / 354	2,400
match	14.9 / 15.3	23.1 / 19.4	174 / 42	2,359
sweep10	77.3 / 74.8	109.5 / 105.4	1,244 / 799	9,635
mixed	30.2 / 30.6	124.8 / 125.2	566 / 556	4,046

- **p50 and p99 barely moved** from WSL2. The code's own speed was measured correctly there.
- **p99.99 fell 5–50×.** On WSL2 every scenario had a floor of about 2.3 µs, the same in every test: that was the virtual machine pausing, not the code. With an isolated core it's gone.
- p99.99 still varies between rounds (e.g. match 174 vs 42 ns): at 1 in 10,000 a few slow samples decide it, so it's reported with both rounds.



Phase 3 latency on bare metal

3.2 Phase 4: the real Nasdaq day through FastBook

`nanobook_replay FILE --throughput --latency --core 8 --csv ...`, twice. Every 8th order message timed (45,389,776 samples).

	Round 1	Round 2	WSL2 (phase 4, commit c13e5bb)
ns per message (throughput)	74.7	75.4	102.3
Parse only	2.7 ns	4.0 ns	3.0 ns
p50, all kinds	85.2	83.9	105
p99	322.8	312.5	429
p99.9	472.1	453.5	—
p99.99	4,001	647	about 11,000
Page faults while timed	0	0	not counted then

Per message kind, round 2 (ns):

Kind	ITCH	p50	p90	p99	p99.99
add	A, F	85.6	131.9	244.7	519
match	E, C	49.2	167.4	300.9	548
cancel	D	64.9	178.2	321.2	634
modify	X	29.8	98.8	237.7	—
replace	U	141.4	269.5	421.6	756

- About a quarter faster than at the end of phase 4: phase 6's kept changes (§4) plus no VM.
- **Round 1's 4.0 µs p99.99 is not trusted.** A few light commands ran on other cores during that round. Round 2 (647 ns) and the eight later full-day runs of the same code (in §4.5 and §4.6, 636–644 ns) agree. **Publish round 2.**

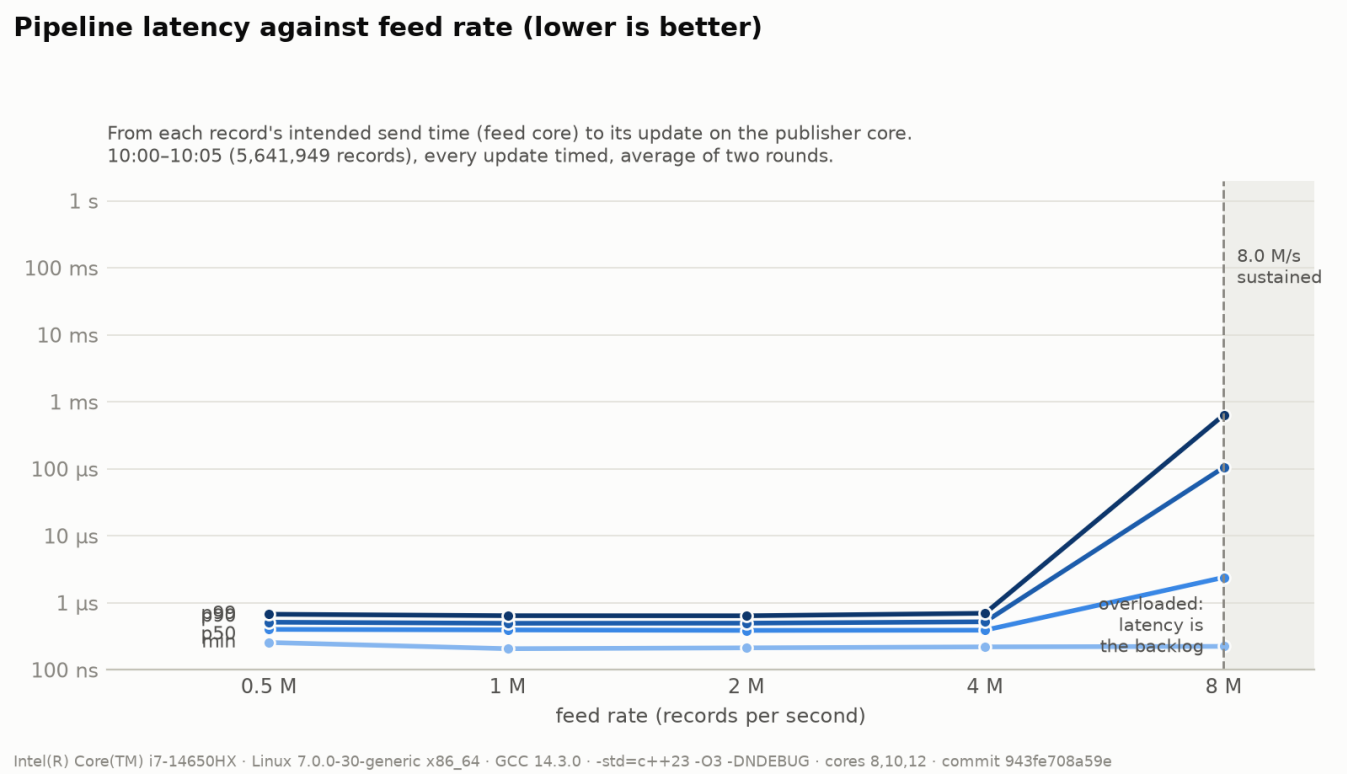
3.3 Phase 5: the three-thread pipeline

`scripts/run_pipeline_session.sh FILE docs/data/bare_metal/phase5 --cores 8,10,12`.

One ring between two cores (`nanobook_ringbench`): round trip p50 175 / 155 ns, p99 204 / 170 ns (WSL2: p50 185 ns). 56–60 ns per record when streaming.

Paced latency, from each record's intended send time to its update leaving the publisher (10:00–10:05, 5,641,949 records, every one timed; round 1 / round 2, ns):

Rate (records/s)	p50	p99	p99.99	WSL2 p99
0.5 M	402 / 403	677 / 678	1,019 / 1,053	65–72 µs
1 M	397 / 392	647 / 641	979 / 956	65–72 µs
2 M	386 / 391	640 / 645	974 / 968	65–72 µs
4 M	394 / 389	703 / 696	4,181 / 1,393	65–72 µs
8 M	overloaded: the backlog grows, latency is the queue			



Phase 5 pipeline latency on bare metal

This answers B4. WSL2's p99 was 65–72 µs at every rate, flat: a sign of pauses that have nothing to do with load. On isolated cores it is **100× lower** and still flat up to 4 M/s. The code was never the problem; the virtual machine was.

Whole day, unpaced: one thread 100.6 ns per message; three threads 124.9 / 125.9 ns (0.80×). Three threads are still slower than one, as on WSL2 (0.69–0.78×). §5.2 shows why.

---

4. Phase 6 re-checks: was each decision right?

Each A/B ran with `scripts/run_ab_session.sh` in ABBA order (before, after, after, before), so a machine that slowly warms up hits both sides equally. Each rule below was fixed **before** the data (optimise design §5). "Noise" is the difference between two runs of the same side.

4.1 B5-E6: huge pages for the books. Confirmed: keep.

E6's commit `43f939e` against its parent, 2 rounds:

	Before (round 1 / 2)	After (round 1 / 2)
ns per message	77.8 / 77.8	72.8 / 72.6
p50	90.1 / 90.1	83.9 / 83.9
p99	347.2 / 346.4	314.2 / 315.4

7% faster in both rounds (noise under 0.3%), p50 and p99 better, the arena "384 MB of it on 2 MiB pages". The worry was that a VM overstates huge pages' gain (a TLB miss there walks two sets of page tables). On bare metal the gain is about the same as on WSL2 (5.6%).

4.2 B5-E2: touching the level arrays at start-up. Fails its rule: revert.

E2's commit `8fd676a` against its parent, 2 rounds:

	Before (round 1 / 2)	After (round 1 / 2)
Page faults while timed	1,650 / 1,650	0 / 0

	Before (round 1 / 2)	After (round 1 / 2)
p99.99	743.6 / 744.5	737.8 / 742.8
p50	90.9 / 90.9	91.8 / 91.4
ns per message	87.0 / 87.0	87.3 / 87.2

E2 removes every page fault, but the rule asks for **p99.99 better in every round beyond the noise**. It moved 6 ns and 2 ns, while two runs of the same code differ by up to 5 ns. On WSL2 a page fault cost much more (the VM handles it twice), which is why E2 looked useful there. **Rule says revert:** the arena stops touching its pages at start-up; the huge pages are then mapped on first use (about 200 faults a day).

4.3 B5-E3: hybrid search near the best price. Confirmed: keep.

E3's commit `352207b` against its parent: 87.6 / 86.9 → 77.9 / 77.8 ns per message (**11% faster**), p50 92.2 → 90.1 ns, p99 360 → 346 ns, cancel p50 89 → 84 ns. Same as WSL2 (10%).

4.4 B5-E1: the feed backs off when ring 1 is full. Confirmed: keep.

E1's commit `7103344` against its parent, whole day:

	Before (round 1 / 2)	After (round 1 / 2)
One thread (control, no ring)	143.5 / 144.0 ns	143.6 / 143.8 ns
Three threads	182.3 / 190.0 ns	156.8 / 155.5 ns
"Ring 1 full" tries	1.36 / 1.39 billion	24.5 / 23.8 million

Three threads **16% faster**; the one-thread control didn't move, so the gain is the ring's.

4.5 B6: E9b, 8-byte ID-table entries. Passes its rule: re-add.

HEAD plus `docs/data/phase6_wsl/e9b_compact/after.diff`, 4 rounds:

	Before (4 rounds)	After (4 rounds)
ns per message	73.0, 72.9, 73.0, 73.0	68.2, 68.2, 68.2, 68.3
p50	84.3 (all 4)	63.7–64.5
p99	315.8–316.2	296.8–297.6

7% faster in every round (noise 0.1 ns), p50 **24% better**, p99 6% better. On WSL2 the same change sat inside the noise and was reverted. The profile said half the replay is in the ID table; halving its entry size fits twice as many entries in each cache line.

4.6 B7: the moving window (E7). Faster, but the slowest messages are slower.

HEAD with `nanobook_replay`'s timing modes switched to `WindowFastBook`, 4 rounds:

	FastBook (4 rounds)	WindowFastBook (4 rounds)
ns per message	72.7–73.0	62.0–62.2
p50	84.3–84.7	64.5–65.7
p99	316.2	317.5–319.5
p99.9	457.6–458.0	477.8–482.4
p99.99	638–643	1,626–4,122

On WSL2 the window's p99 was 18–26% worse; on bare metal p99 is equal and the damage has moved further out, to the slowest 1 in 10,000. The window is 15% faster on average.

---

5. What the hardware counters say

WSL2 can't read the CPU's performance counters; bare metal can. All counters are from the performance core the code ran on (`cpu_core`), with `perf stat`.

5.1 B1: where a real-day message's time goes

`perf stat -I 1000` on `nanobook_replay --throughput`; the one-second intervals separate the counting pass, the parse-only pass and the book pass. Book pass, per message, two runs each:

Per message	FastBook	WindowFastBook (B7)
Cycles	377–378	322–325
Instructions per cycle	0.88	0.92–0.93
L1 data misses	3.85	3.94–4.00
L2 misses	1.87	1.89–1.90
L3 misses (trips to main memory)	**0.52**	**0.56–0.57**
TLB misses caught by the second-level TLB	3.7	3.2
Page walks	**0.05**	0.04

- **Main memory dominates.** About one message in two waits for DRAM (around 80–100 ns each). That fits the low instructions per cycle (0.88: the CPU is mostly waiting).
- **Address translation is solved.** With huge pages, a page walk happens once in 20 messages. The counting pass, which doesn't use the arena, does 1.2 walks per message: that is what the books would cost without E6.
- **B7's hypothesis is supported but not proven.** The window makes about 9% more trips to main memory per message, which fits "a sparse window spreads levels over more lines" and could explain its slower tail. Averages can't prove what causes a p99.99, though.

5.2 B2: cache lines moving between cores in the pipeline

`perf c2c` (the usual tool) found no cross-core "modified line" loads in user code on this CPU, in two attempts: its decoding doesn't cover this hybrid client chip. So the session counted them directly: `mem_load_l3_hit_retired.xsnp_hitm` per core, every 2 s, over a whole-day run.

Phase of the run	Feed (CPU 8)	Book (CPU 10)	Publisher (CPU 12)
One thread (all on CPU 10)	—	0	—
Three threads, per message	about 0.09	about 0.11	**about 2.5**

- **Ring 1 is quiet.** E1's back-off works: the feed and book cores rarely steal each other's lines.
- **The publisher is the hot spot.** It spins on ring 2, and each look at a line the book core just wrote moves that line across. This is what E1d (batching ring 2's head) attacked: it made three threads 7–12% faster but the paced p90 worse, so it wasn't kept.

5.3 B3: the `match` puzzle is gone

On WSL2, FastBook's `match` took 9.3 ns in the harness's throughput mode but 13.1 ns in Google Benchmark (29% apart), unexplained. On bare metal: **8.7 ns in both**. It was the virtual machine.

6. Decisions for the user

#	Question	What the rule or data says
1	**E2**: stop pre-touching the arena?	The rule fixed on 2026-10-21 says **revert**: no measurable p99.99 gain
2	**E9b**: re-add 8-byte ID entries?	The rule fixed on 2026-10-22 says **re-add**: 7% faster, p50 24% better, 4 of 4 rounds
3	**E7**: adopt the window?	No rule decides it: 15% faster and p99 equal, but p99.99 2.5–6× worse

After the decisions: write the perf-log entries, update optimise design §5 and the README status, mark phases 3–6 complete, and commit the data. The two `perf.data` files from B2 (about 17 MB, which found nothing) can stay out of git.

---

7. Data and how to reproduce

Folder in docs/data/bare_metal/	What
`phase3/`	Latency session: two rounds, throughput mode, Google Benchmark, charts
`phase4/`	Two full-day replays (`.txt`, `.csv`, histograms)
`phase5/`	Ring benchmark, whole-day throughput, paced latency at five rates, two rounds each
`phase6/b5_e6`, `b5_e2`, `b5_e3`, `b5_e1`, `b6_e9b`, `b7_e7`	The A/B sessions (`progress.log`, every run's output, `after.diff` where there is one)
`phase6/perf_b1_b7/`	Interval counters for B1 and B7; `b1_per_message.txt`
`phase6/snoop_b2/`	Cross-core counters for B2
`phase6/perf_b3/`	Counters around `match` in both tools

To repeat on the same stick: boot with the three options of §2.2, mount the Windows drive read-only, set the `performance` governor, then run the commands quoted in each section.